

SHELL I/O REDIRECTIONS

Most Unix system commands take input from your terminal and send the resulting output back to your terminal. A command normally reads its input from the standard input, which happens to be your terminal by default. Similarly, a command normally writes its output to standard output, which is again your terminal by default.

Output Redirection

The output from a command normally intended for standard output can be easily diverted to a file instead. This capability is known as output redirection.

If the notation `> file` is appended to any command that normally writes its output to standard output, the output of that command will be written to file instead of your terminal.

Check the following **who** command which redirects the complete output of the command in the users file.

```
$ who > users
```

Notice that no output appears at the terminal. This is because the output has been redirected from the default standard output device (the terminal) into the specified file. You can check the users file for the complete content –

```
$ cat users
oko      tty01  Sep 12 07:30
ai       tty15  Sep 12 13:32
ruth     tty21  Sep 12 10:10
pat      tty24  Sep 12 13:07
steve    tty25  Sep 12 13:03
$
```

If a command has its output redirected to a file and the file already contains some data, that data will be lost. Consider the following example –

```
$ echo line 1 > users
$ cat users
line 1
$
```

You can use `>>` operator to append the output in an existing file as follows –

```
$ echo line 2 >> users
```

```
$ cat users  
line 1  
line 2  
$
```

Input Redirection

Just as the output of a command can be redirected to a file, so can the input of a command be redirected from a file. As the **greater-than character** `>` is used for output redirection, the **less-than character** `<` is used to redirect the input of a command.

The commands that normally take their input from the standard input can have their input redirected from a file in this manner. For example, to count the number of lines in the file *users* generated above, you can execute the command as follows –

```
$ wc -l users  
2 users  
$
```

Upon execution, you will receive the following output. You can count the number of lines in the file by redirecting the standard input of the **wc** command from the file *users* –

```
$ wc -l < users  
2  
$
```

Note that there is a difference in the output produced by the two forms of the **wc** command. In the first case, the name of the file *users* is listed with the line count; in the second case, it is not.

In the first case, **wc** knows that it is reading its input from the file *users*. In the second case, it only knows that it is reading its input from standard input so it does not display file name.